

COMP809 Data Mining and Machine Learning

Patricio Maturana-Russel

`p.maturana.russel@aut.ac.nz`

*Department of Mathematical Sciences and Computer Science and Software Engineering
Departments, Auckland University of Technology, Auckland, New Zealand*

Semester 1, 2024

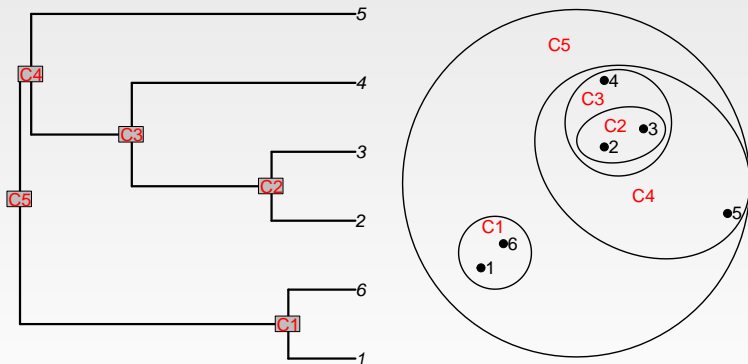


Contents

- Hierarchical Clustering
- Agglomerative
- Case study

Hierarchical Clustering

- It produces a set of nested clusters organized as a hierarchical tree.
- It can be visualized as a dendrogram—a tree like diagram that records the sequences of merges or splits.



Hierarchical Clustering

Two main types:

- Agglomerative (bottom-up):
 - Start with the **points as individual clusters**.
 - At each step, merge the closest pair of clusters until only one cluster (or K clusters) left.
- Divisive (top-down):
 - Start with one, **all-inclusive cluster**.
 - At each step, split a cluster until each cluster contains a point (or there are K clusters)

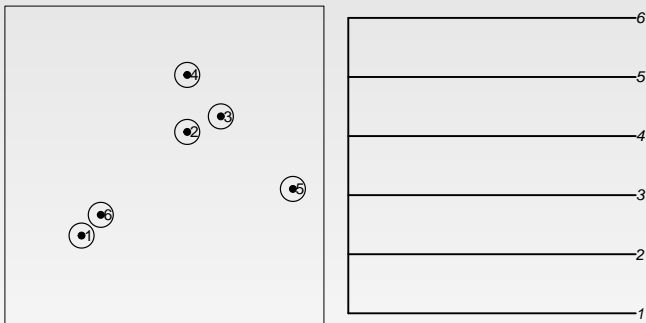
Traditional hierarchical algorithms use a similarity or distance matrix, and merge or split one cluster at a time.

Agglomerative

- Most popular hierarchical clustering technique (simpler mathematically).
- Basic algorithm is straightforward:
 - Compute the proximity matrix
 - Let each data point be a cluster
 - **Repeat**
 - Merge the two closest clusters
 - Update the proximity matrix
 - **Until** only a single cluster remains
- Key operation is the computation of the proximity of two clusters.
 - Different approaches to defining the distance between clusters distinguish the different algorithms.

Agglomerative

Each data point is a cluster.

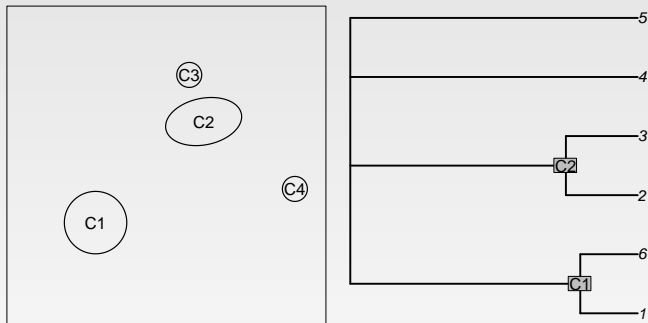


Proximity Matrix

	1	2	3	4	5	6
1						
2						
3						
4						
5						
6						

Agglomerative

After a few iterations:

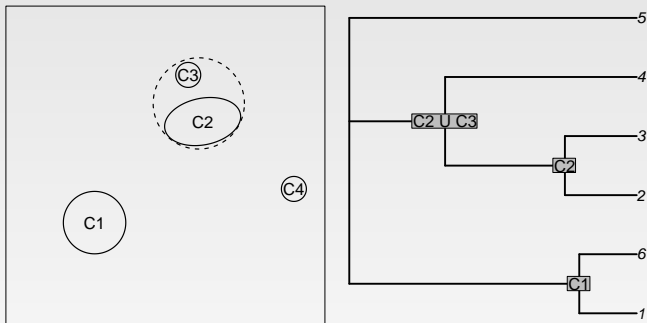


Proximity Matrix

	C1	C2	C3	C4
C1				
C2				
C3				
C4				

Agglomerative

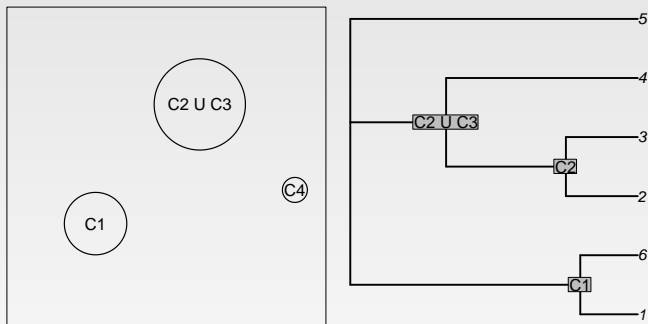
Identify and merge the closest cluster



Proximity Matrix

	C1	C2	C3	C4
C1				
C2				
C3				
C4				

Agglomerative

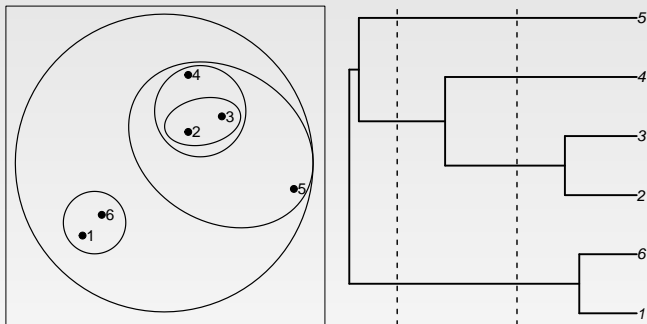


Proximity Matrix

	C1	C2 U C3	C4
C1			
C2 U C3			
C4			

Continue until one cluster remain.

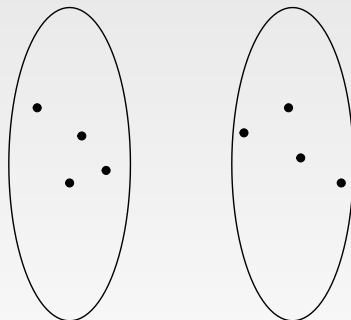
Agglomerative



Clustering obtained by cutting the dendrogram at a desired level: each connected component forms a cluster.

Inter-cluster similarity

- MIN (single-link)
- MAX (complete-link)
- Group Average (average-link)
- Distance Between Centroids (centroid)
- Ward

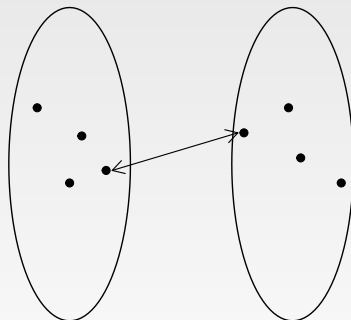


Inter-cluster similarity

- **MIN (single-link)**
- MAX (complete-link)
- Group Average (average-link)
- Distance Between Centroids (centroid)
- Ward

Limitations:

- Sensitive to noise and outliers.
- It produces long, elongated clusters.



Inter-cluster similarity

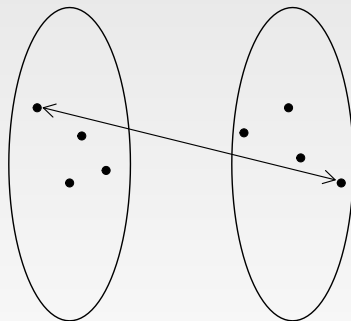
- MIN (single-link)
- **MAX (complete-link)**
- Group Average (average-link)
- Distance Between Centroids (centroid)
- Ward

Strength:

- More balanced clusters (with equal diameter).
- Less susceptible to noise.

Limitations:

- Tends to break large clusters.
- All clusters tend to have the same diameter—small clusters are merged with larger ones.



Inter-cluster similarity

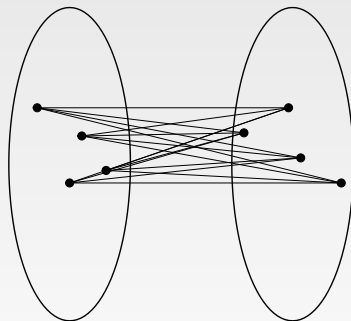
- MIN (single-link)
- MAX (complete-link)
- **Group Average (average-link)**
- Distance Between Centroids (centroid)
- Ward

Strength:

- Less susceptible to noise and outliers

Limitations:

- Biased towards globular clusters

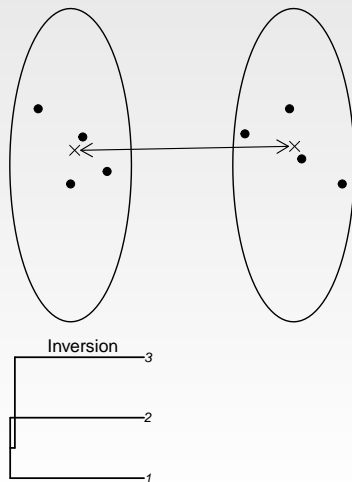


Inter-cluster similarity

- MIN (single-link)
- MAX (complete-link)
- Group Average (average-link)
- **Distance Between Centroids (centroid)**
- Ward

Limitations:

- It is not monotonic: smaller clusters can potentially be more similar to the new larger cluster than to their individual clusters causing an inversion in the dendrogram. It contradicts the fundamental assumption that small clusters are more coherent than large clusters. It does not arise in the other methods.



Inter-cluster similarity

- MIN (single-link)
- MAX (complete-link)
- Group Average (average-link)
- Distance Between Centroids (centroid)
- **Ward**

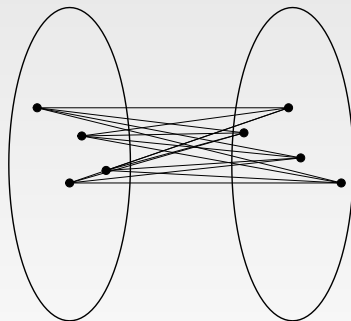
It minimizes the variance of the clusters being merged.

Strength:

- Less susceptible to noise and outliers.

Limitations:

- Biased towards globular clusters.



Inter-cluster similarity

Ward's distance between clusters C_i and C_j is defined as:

$$D_w(C_i, C_j) = \sum_{x \in C_i} (x - r_i)^2 + \sum_{x \in C_j} (x - r_j)^2 - \sum_{x \in C_{ij}} (x - r_{ij})^2,$$

where r_i , r_j , and r_{ij} are the centroids of C_i , C_j , and $C_{ij} = C_i \cup C_j$, respectively.

It is the difference between the sum of the squared errors of clusters C_i and C_j , and the squared error of the merging of both clusters.

Ward's method is the hierarchical analogue of K-means. It can be used to initialize K-means.

Hierarchical Clustering

Advantages:

- It shows all the possible linkages between clusters, which helps to understand the data.
- No need of assuming the number of clusters.
 - Any desired number of clusters can be obtained by 'cutting' the dendrogram at the proper level.
- They may correspond to meaningful taxonomies.
 - Example in biological sciences (e.g., animal kingdom, phylogeny reconstruction, etc).

Hierarchical Clustering

Disadvantages:

- In large data sets
 - It can be difficult to interpret.
 - It can be computationally expensive.
- Once a decision is made to combine two clusters, it cannot be undone.
- No objective function is directly minimized.
- Different schemes have problems with one or more of the following:
 - Sensitivity to noise and outliers.
 - Difficulty handling different sized clusters and convex shapes.
 - Breaking large clusters.

In Python

First, we produce the dendrogram to determine the number of clusters

```
>>> from scipy.cluster.hierarchy import dendrogram, linkage;  
>>> dendrogram(linkage(data, method="ward"), orientation = "top",  
                labels = None);
```

Find more information on www.docs.scipy.org

Then can get the clusters as follows:

```
>>> from sklearn.cluster import AgglomerativeClustering  
>>> model = AgglomerativeClustering(n_clusters=2, linkage="ward")  
>>> model.fit(data)
```

Find more information on www.scikit-learn.org

Case study

The dataset has been compiled from the United Nations Demographic Yearbook 1990 (United Nations publications) and has the following variables: *birth rate*, *death rate*, *infant death rate*, and *country*.

Can these variables be used to categorize these countries?

```
>>> import pandas as pd
>>> data = pd.read_csv("poverty.csv");
>>> print(data);
```

	Birth	Death	InfantDeath	Country
0	24.7	5.7	30.8	Albania
1	13.4	11.7	11.3	Czechoslovakia
2	11.6	13.4	14.8	Hungary
3	13.6	10.7	26.9	Romania
4	17.7	10.0	23.0	USSR
..
92	50.1	20.2	132.0	Somalia
93	44.6	15.8	108.0	Sudan
94	31.1	7.3	52.0	Tunisia
95	50.5	14.0	106.0	Tanzania
96	51.1	13.7	80.0	Zambia

Case study

```
>>> data[["Birth", "Death", "InfantDeath"]].describe();
```

	Birth	Death	InfantDeath
count	97.000000	97.000000	97.000000
mean	29.229897	10.836082	54.901031
std	13.546695	4.647495	45.992584
min	9.700000	2.200000	4.500000
25%	14.500000	7.800000	13.100000
50%	29.000000	9.500000	43.000000
75%	42.200000	12.500000	83.000000
max	52.200000	25.000000	181.600000

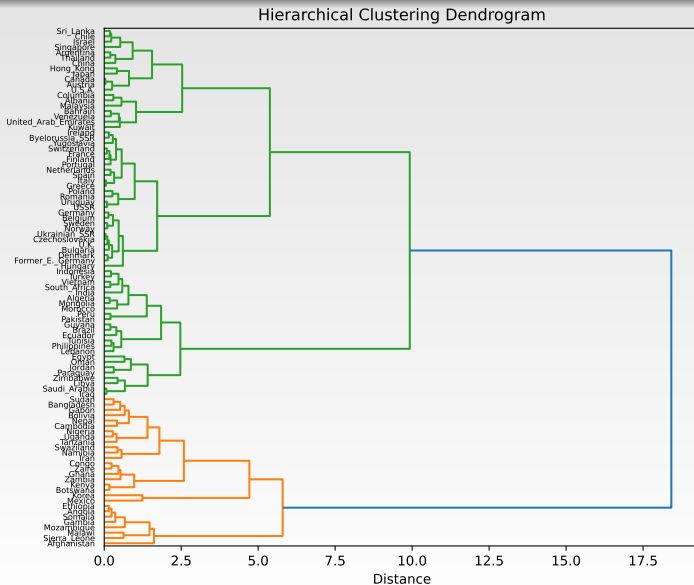
The SDs are quite different. The data will be standardized.

```
>>> from sklearn.preprocessing import StandardScaler
>>> X      = data.iloc[:, [0,1,2]];
>>> scaler = StandardScaler();
>>> fitted = scaler.fit(X)
>>> X_std  = pd.DataFrame(fitted.transform(X));
```

We can generate the dendrogram as follows:

```
>>> from scipy.cluster.hierarchy import dendrogram, linkage;
>>> dendrogram(linkage(X_std, method="ward"), orientation = "right",
               labels = data.Country.tolist());
```

Case study



Case study

The number of clusters can be inferred from the dendrogram by drawing a vertical line on it. This should be where we find the biggest distances. In our example, it could be between approximately

- 10 and 18, generating 2 clusters; or
- 5.5 and 10, generating 3 clusters.

Case study

We can add the cluster label to the original dataset for further analyses as follows:

```
>>> from sklearn.cluster import AgglomerativeClustering
>>> model = AgglomerativeClustering(n_clusters=3, linkage="ward",
                                   compute_distances = True);
>>> model.fit(X_std);

>>> data["Cluster"] = pd.DataFrame(model.labels_);
>>> print(data);
```

	Birth	Death	InfantDeath	Country	Cluster
0	24.7	5.7	30.8	Albania	1
1	13.4	11.7	11.3	Czechoslovakia	1
2	11.6	13.4	14.8	Hungary	1
3	13.6	10.7	26.9	Romania	1
4	17.7	10.0	23.0	USSR	1
..
92	50.1	20.2	132.0	Somalia	0
93	44.6	15.8	108.0	Sudan	0
94	31.1	7.3	52.0	Tunisia	2
95	50.5	14.0	106.0	Tanzania	0
96	51.1	13.7	80.0	Zambia	0

[97 rows x 5 columns]

Case study

To describe the clusters, we need to study its components:

```
>>> print("Cluster 1:\n", list(data["Country"][(data["Cluster"]==0)]));
Cluster 1:
['Bolivia', 'Mexico', 'Afghanistan', 'Iran', 'Bangladesh', 'Korea', 'Botswana',
 'Gabon', 'Ghana', 'Namibia', 'Sierra_Leone', 'Swaziland', 'Uganda', 'Zaire',
 'Cambodia', 'Nepal', 'Angola', 'Congo', 'Ethiopia', 'Gambia', 'Kenya', 'Malawi',
 'Mozambique', 'Nigeria', 'Somalia', 'Sudan', 'Tanzania', 'Zambia']

>>> print("Cluster 2:\n", list(data["Country"][(data["Cluster"]==1)]));
Cluster 2:
['Albania', 'Czechoslovakia', 'Hungary', 'Romania', 'USSR', 'Ukrainian_SSR',
 'Chile', 'Uruguay', 'Finland', 'France', 'Greece', 'Italy', 'Norway', 'Spain',
 'Switzerland', 'Austria', 'Canada', 'Israel', 'Kuwait', 'China', 'Singapore',
 'Thailand', 'Bulgaria', 'Former_E_Germany', 'Poland', 'Yugoslavia',
 'Byelorussia_SSR', 'Argentina', 'Columbia', 'Venezuela', 'Belgium', 'Denmark',
 'Germany', 'Ireland', 'Netherlands', 'Portugal', 'Sweden', 'U.K.', 'Japan',
 'U.S.A.', 'Bahrain', 'United_Arab_Emirates', 'Hong_Kong', 'Malaysia', 'Sri_Lanka']

>>> print("Cluster 3:\n", list(data["Country"][(data["Cluster"]==2)]));
Cluster 3:
['Ecuador', 'Paraguay', 'Oman', 'Turkey', 'India', 'Mongolia', 'Pakistan',
 'Algeria', 'Egypt', 'Libya', 'Morocco', 'South_Africa', 'Zimbabwe', 'Brazil',
 'Guyana', 'Peru', 'Iraq', 'Jordan', 'Lebanon', 'Saudi_Arabia', 'Indonesia',
 'Philippines', 'Vietnam', 'Tunisia']
```

Case study

The classification of a new observation is not possible through agglomerative algorithms because they do not partition the input space, they just connects some objects during the clustering.

A potential solution could be the use of a supervised machine learning technique, using the output from the agglomerative algorithms as labels, and then classify the new observations accordingly. For instance, we could use the k-nearest neighbour algorithm as follows to classify the following hypothetical countries:

	Birth	Death	InfantDeath
Country A	10	3	5
Country B	29	11	55
Country C	52	25	180

```
>>> from sklearn.neighbors import KNeighborsClassifier
>>> new_data = pd.DataFrame([[10,3,5], [29, 11, 55], [52, 25, 180]],
                           columns=["Birth", "Death", "InfantDeath"])
>>> new_data_std = pd.DataFrame(fitted.transform(new_data))
>>> knn = KNeighborsClassifier(n_neighbors=1)
>>> knn.fit(X_std, model.labels_)
>>> print(knn.predict(new_data)+1)
[2 1 3]
```

End