

COMP824 2023 Week 5

Transforming Data

Dr Sarah Marshall
sarah.marshall@aut.ac.nz WZ Level 9

Department of Mathematical Sciences
Auckland University of Technology



Overview

The Process of Analytics

Tibbles

Transforming Data `dplyr`

Workflow: R scripts

Reading

Chapter 5, 6, 10 Wickham and Grolemund (2020), R for Data Science

<https://r4ds.had.co.nz/>

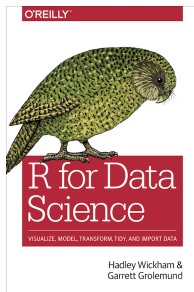



Figure 1: <http://r4ds.had.co.nz/>

The Process of Analytics



Learning objectives

- Understand the difference between data frames and tibbles
- Know how to filter, arrange, select, mutate and summarise data with `dplyr`
- Appreciate the beauty and usefulness of “the pipe” `%>%`
- Understand the key principles of using R scripts

Tibbles

Definition: Tibbles are “optionated data frames that make working in the tidyverse a little easier”. Wickham and Grolemund (2020)



Figure 3: [https://https://www.tidyverse.org/](https://www.tidyverse.org/)

Example: Data frames

```
head(iris)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

Example: Tibbles

```
as_tibble(iris)
```

```
# A tibble: 150 x 5
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
	<dbl>	<dbl>	<dbl>	<dbl>	<fct>
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa

```
# ... with 147 more rows
```

```
class(as_tibble(iris))
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```


Creating Tibbles

- Coerce a data frame `as_tibble`
- Create a new tibble `tibble`

Creating Tibbles

```
tibble(  
  x = 1:3,  
  y = 1,  
  z = x ^ 2 + y,  
  `x squared` = x^2  
)
```

A tibble: 3 x 4

	x	y	z	`x squared`
	<int>	<dbl>	<dbl>	<dbl>
1	1	1	2	1
2	2	1	5	4
3	3	1	10	9

Use backticks ' to define and access non-traditional column names (e.g. spaces, starting with numbers etc.)

Tribbles

Transposed tibbles – for easier data entry

```
tribble(  
  ~x, ~y, ~z,  
  #--|--|----  
  "a", 2, 3.6,  
  "b", 1, 8.5  
)
```

```
# A tibble: 2 x 3  
  x          y      z  
  <chr> <dbl> <dbl>  
1 a           2    3.6  
2 b           1    8.5
```

Tibbles vs Data Frames: Default Printing Options

- Tibbles: max 10 rows, limited columns, variable type
- Data frame: (almost) all rows, all columns

Tibbles vs Data Frames: Customizing Printing Options 1

- Tibble (local)

```
nycflights13::flights %>%  
  print(n = 10, width = Inf)
```

- Tibble (global)

```
options(pillar.print_max = n,  
        pillar.print_min = m,  
        tibble.width = Inf)
```

Tibbles vs Data Frames: Customing Printing Options 2

- Tibble (view)

```
nycflights13::flights %>% View()
```

- Data frame

```
head(iris, 10)  
tail(iris)  
?print.data.frame
```

Tibbles vs Data Frames: Subsetting

- `$` extract a variable by name
- `[[]]` extract by a variable name or position

Tibbles vs Data Frames: Subsetting (Example)

```
tb <- tibble(x = 1:5,  
             y = 11:15)
```

```
# Extract by name  
tb$x; tb[["x"]]
```

```
[1] 1 2 3 4 5
```

```
[1] 1 2 3 4 5
```

```
# Extract by position  
tb[[2]]
```

```
[1] 11 12 13 14 15
```

```
# Extract using a pipe - note the dot  
tb %>% .$x
```


Tibbles to data frames

```
as.data.frame(tb)
```

```
   x  y  
1  1 11  
2  2 12  
3  3 13  
4  4 14  
5  5 15
```

```
class(tb)
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

```
class(as.data.frame(tb))
```

```
[1] "data.frame"
```

Transforming Data `dplyr`

Description: `dplyr` is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges

<https://dplyr.tidyverse.org/>



Figure 4: <https://https://www.tidyverse.org/>

Transforming data with dplyr

5 key functions

- `filter()` – pick observations based on criteria
- `arrange()` – reorder the rows based on a column value
- `select()` – select specified columns by name
- `mutate()` – create new variables from existing ones
- `summarise()` – create summaries

Also useful:

- `group_by()` – apply functions by group
- `rename()` – renames variables

2013 NYC Flights Data

```
library(nycflights13)  
library(tidyverse)
```

```
?flights
```

2013 NYC Flights Data

flights

```
# A tibble: 336,776 x 19
  year month   day dep_time sched~1 dep_d~2 arr_t~3 sched~4
  <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>
1  2013     1     1     517     515     2     830     819
2  2013     1     1     533     529     4     850     830
3  2013     1     1     542     540     2     923     850
# ... with 336,773 more rows, 11 more variables:
#   arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>, and abbreviated
#   variable names 1: sched_dep_time, 2: dep_delay,
#   3: arr_time, 4: sched_arr_time
```

Filter

filter – pick observations based on criteria

```
(jan2 <- filter(flights, month == 1, day == 2))
```

```
# A tibble: 943 x 19
```

	year	month	day	dep_time	sched~1	dep_d~2	arr_t~3	sched~4
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	2	42	2359	43	518	442
2	2013	1	2	126	2250	156	233	2359
3	2013	1	2	458	500	-2	703	650

```
# ... with 940 more rows, 11 more variables:
```

```
#   arr_delay <dbl>, carrier <chr>, flight <int>,
```

```
#   tailnum <chr>, origin <chr>, dest <chr>,
```

```
#   air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
#   minute <dbl>, time_hour <dtm>, and abbreviated
```

```
#   variable names 1: sched_dep_time, 2: dep_delay,
```

```
#   3: arr_time, 4: sched_arr_time
```

Filter - with logic

More examples:

```
on_schedule <- filter(flights, arr_delay <= 20, dep_delay <= 20)

on_time_departure <- filter(flights, between(dep_delay, -1, 1))

# The following two statements give the same results
nov_dec <- filter(flights, month == 11 | month == 12)
nov_dec_v2 <- filter(flights, month %in% c(11, 12))
```

Arrange

arrange - reorder the rows based on a column value

```
arrange(flights, month, day, sched_dep_time)
```

```
# A tibble: 336,776 x 19
```

	year	month	day	dep_time	sched~1	dep_d~2	arr_t~3	sched~4
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	517	515	2	830	819
2	2013	1	1	533	529	4	850	830
3	2013	1	1	542	540	2	923	850

```
# ... with 336,773 more rows, 11 more variables:  
#   arr_delay <dbl>, carrier <chr>, flight <int>,  
#   tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>,  
#   minute <dbl>, time_hour <dtm>, and abbreviated  
#   variable names 1: sched_dep_time, 2: dep_delay,  
#   3: arr_time, 4: sched_arr_time
```


Arrange - Further examples

Question: Which flights had the longest departure delays?

```
arrange(flights, dep_delay)
arrange(flights, desc(dep_delay)) #Descending order
```

```
arrange(flights, -dep_delay)      #Descending order
```

```
# A tibble: 336,776 x 19
```

	year	month	day	dep_time	sched~1	dep_d~2	arr_t~3	sched~4
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	9	641	900	1301	1242	1530
2	2013	6	15	1432	1935	1137	1607	2120
3	2013	1	10	1121	1635	1126	1239	1810

```
# ... with 336,773 more rows, 11 more variables:
```

```
#   arr_delay <dbl>, carrier <chr>, flight <int>,
```

```
#   tailnum <chr>, origin <chr>, dest <chr>,
```

```
#   air_time <dbl>, distance <dbl>, hour <dbl>,
```

Select

select - select specified columns by name

```
select(flights, year, month, day)
```

```
# A tibble: 336,776 x 3
```

```
  year month   day
```

```
  <int> <int> <int>
```

```
1  2013     1     1
```

```
2  2013     1     1
```

```
3  2013     1     1
```

```
# ... with 336,773 more rows
```

Select - Further examples 1

```
# Select all columns from year to day  
select(flights, year:day)
```

```
#Select all columns except year:day  
select(flights, -(year:day))
```

```
# Select columns related to departures + tailnum  
select(flights, starts_with("dep"), tailnum)  
select(flights, contains("dep"))
```

```
# Move time_hour and air_time to start then  
# include remaining columns  
select(flights, time_hour, air_time, everything())
```

Select - Further examples 2

```
(flights_sml <- select(flights,  
  year:day,  
  ends_with("delay"),  
  distance,  
  air_time  
))
```

A tibble: 336,776 x 7

	year	month	day	dep_delay	arr_delay	distance	air_time
	<int>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
1	2013	1	1	2	11	1400	227
2	2013	1	1	4	20	1416	227
3	2013	1	1	2	33	1089	160

... with 336,773 more rows

Mutate - create new variables

```
mutate(flights_sml,  
       gain = arr_delay - dep_delay,  
       hours = air_time / 60,  
       gain_per_hour = gain / hours  
)
```

A tibble: 336,776 x 10

	year	month	day	dep_delay	arr_de~1	dista~2	air_t~3	gain
	<int>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	2013	1	1	2	11	1400	227	9
2	2013	1	1	4	20	1416	227	16
3	2013	1	1	2	33	1089	160	31

... with 336,773 more rows, 2 more variables:

hours <dbl>, gain_per_hour <dbl>, and abbreviated

variable names 1: arr_delay, 2: distance, 3: air_time

Transmute - only keep new variables

```
#Keep only new variables
transmute(flights,
          gain = arr_delay - dep_delay,
          hours = air_time / 60,
          gain_per_hour = gain / hours
)
```

```
# A tibble: 336,776 x 3
```

	gain	hours	gain_per_hour
	<dbl>	<dbl>	<dbl>
1	9	3.78	2.38
2	16	3.78	4.23
3	31	2.67	11.6

```
# ... with 336,773 more rows
```

Mutate + Modular arithmetic

```
transmute(flights,  
  dep_time,  
  hour = dep_time %/% 100,  
  minute = dep_time %% 100  
)
```

A tibble: 336,776 x 3

	dep_time	hour	minute
	<int>	<dbl>	<dbl>
1	517	5	17
2	533	5	33
3	542	5	42

... with 336,773 more rows

Summarise and Group By

`group_by` - create groups within dataframe

`summarise` - summarise by group

Discussion: Which month has the greatest average departure delay?

Summarise and Group By: Departure Delays By Month

```
by_month <- group_by(flights, month)
summarise(by_month, delay = mean(dep_delay, na.rm = TRUE))
```

```
# A tibble: 12 x 2
  month delay
  <int> <dbl>
1     1  10.0
2     2  10.8
3     3  13.2
# ... with 9 more rows
```

Summarise and Group By: Departure Delays By Month 2

```
by_month <- group_by(flights, month)
delay_by_month <- summarise(by_month,
                             delay = mean(dep_delay, na.rm = TRUE))
arrange(delay_by_month, -delay)
```

```
# A tibble: 12 x 2
  month delay
  <int> <dbl>
1     7  21.7
2     6  20.8
3    12  16.6
# ... with 9 more rows
```

Summarise and Group By: Arrival Delays by Month

```
summarise(by_month,  
  avg_arr_delay1 = mean(arr_delay, na.rm = TRUE),  
  avg_arr_delay2 = mean(arr_delay[arr_delay > 0], na.rm = TRUE)  
  # the average positive delay  
)
```

A tibble: 12 x 3

	month	avg_arr_delay1	avg_arr_delay2
	<int>	<dbl>	<dbl>
1	1	6.13	34.5
2	2	5.61	33.7
3	3	5.81	40.6
# ...	with 9 more rows		

Summarise and Group By: Number of flights per day

```
daily <- group_by(flights, year, month, day)
(per_day <- summarise(daily, flights = n()))
```

```
# A tibble: 365 x 4
# Groups:   year, month [12]
   year month   day flights
  <int> <int> <int>    <int>
1  2013     1     1     842
2  2013     1     2     943
3  2013     1     3     914
# ... with 362 more rows
```

Summarise and Group By: Number of flights per month

Each summarise removes a layer of the grouping.

```
(per_month <- summarise(per_day, flights = sum(flights)))
```

```
# A tibble: 12 x 3
# Groups:   year [1]
   year month flights
  <int> <int>   <int>
1  2013     1   27004
2  2013     2   24951
3  2013     3   28834
# ... with 9 more rows
```

Summarise and Group By: Number of flights per year

```
(per_year <- summarise(per_month, flights = sum(flights)))
```

```
# A tibble: 1 x 2  
  year flights  
  <int>   <int>  
1  2013  336776
```

The pipe %>%



Figure 5: <https://https://www.tidyverse.org/>

The pipe %>%

The pipe operator %>% makes analysis much easier.

It “pipes” the results of a function into the first argument of the next function.

```
x <- 1:4
```

```
# Without a pipe
```

```
mean(x^2)
```

```
[1] 7.5
```

```
# With a pipe
```

```
x^2 %>% mean()
```

```
[1] 7.5
```


The pipe %>%

```
# Without a pipe
by_month <- group_by(flights, month)
delay_by_month <- summarise(by_month,
                             delay = mean(dep_delay, na.rm = TRUE))
arrange(delay_by_month, -delay)

# With a pipe
flights %>%
  group_by(month) %>%
  summarise(delay = mean(dep_delay, na.rm = TRUE)) %>%
  arrange(-delay)
```

The new pipe |>

In R 4.1 and later there is a new pipe operator built-in to Base R.

It works a similar way to the `maggritr` pipe, but there are some differences

```
x %>% mean()  
x %>% mean  
x |> mean()  
x |> mean    # error: the new pipe needs brackets
```

For more info: <https://www.infoworld.com/article/3621369/use-the-new-r-pipe-built-into-r-41.html>



Workflow: R scripts

R scripts

R scripts are files containing R commands

The screenshot displays the RStudio environment with the following components:

- Source Editor:** Contains an R script named `busy_flights.R` with the following code:

```
1 # Load packages
2 library(dplyr)
3 library(nycflights13)
4
5 # Find busy days
6 flights %>%
7   group_by(year, month, day) %>%
8   summarise(numflights = n()) %>%
9   arrange(desc(numflights))
10
```
- Console:** Shows the output of running the script:

```
> library(dplyr)
> library(nycflights13)
> flights %>%
+   group_by(year, month, day) %>%
+   summarise(numflights = n()) %>%
+   arrange(desc(numflights))
# A tibble: 365 x 4
# Groups:   year, month [12]
   year month   day numflights
  <int> <int> <int>    <int>
1  2013    11    27        1014
2  2013     7    11        1006
3  2013     7     8         1004
4  2013     7    10         1004
5  2013    12     2         1004
6  2013     7    18         1003
7  2013     7    25         1003
8  2013     7    12         1002
9  2013     7     9         1001
10 2013     7    17         1001
# ... with 355 more rows
>
```
- Environment:** Shows the Global Environment, which is currently empty.
- Files, Plots, Packages, Help, Viewer:** The Help pane is active, displaying the R documentation for the `arrange` function from the `dplyr` package.

R Documentation: `arrange` (dplyr)

Arrange rows by variables

Description

Use `desc()` to sort a variable in descending order.

Usage

```
arrange(.data, ...)
```

S3 method for class 'grouped_df'

```
arrange(.data, ..., by_group = FALSE)
```

R scripts

Good practice

- Load packages at the start
- Use comments
- Store code not results
- Never save or restore .Rdata into workspace (Tools/Global Options/General)
- Use relative not absolute paths
- If sharing files – including `install.packages()` or `setwd()` is “anti-social”

Keyboard shortcuts

- `Cmd/Ctrl + Enter`: Run current R expression
- `Cmd/Ctrl + Shift + S`: Run/Source complete R script
- `Cmd/Ctrl + Shift + F10`: Restart RStudio

Learning objectives

- Understand the difference between data frames and tibbles
- Know how to filter, arrange, select, mutate and summarise data with `dplyr`
- Appreciate the beauty and usefulness of “the pipe” `%>%`
- Understand the key principles of using R scripts



References

Wickham, Hadley, and Garrett Grolemund. 2020. *R for Data Science: Import, Tidy, Transform, Visualize, and Model Data*.