

COMP824 2023 Week 9

Relational Data

Department of Mathematical Sciences
Auckland University of Technology



Overview

Relational data

Keys

Joins

Visualising Geographic Data

Reading

Chapter 13 Wickham and Grolemund (2020), R for Data Science

<https://r4ds.had.co.nz/>

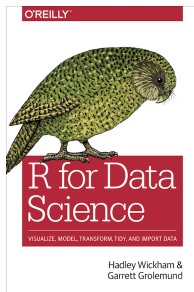



Figure 1: <http://r4ds.had.co.nz/>

The Process of Analytics





Learning objectives

- Recognise relational data
- Understand the main types of mutating and filtering joins
- Join datasets using appropriate `tidyverse` join functions

Relational data

Multiple tables of related data = **relational data**

Example

- Each flight has an airline.
- Each airline has multiple flights
- The tibbles `airlines` and `flights` are related.

Flights data

```
nycflights13::flights
```

```
# A tibble: 336,776 x 19
  year month   day dep_time sched~1 dep_d~2 arr_t~3 sched~4
  <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>
1  2013     1     1     517     515       2     830     819
2  2013     1     1     533     529       4     850     830
3  2013     1     1     542     540       2     923     850
# ... with 336,773 more rows, 11 more variables:
#   arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>, and abbreviated
#   variable names 1: sched_dep_time, 2: dep_delay,
#   3: arr_time, 4: sched_arr_time
```

Airline data

```
nycflights13::airlines
```

```
# A tibble: 16 x 2
  carrier name
  <chr>    <chr>
1 9E      Endeavor Air Inc.
2 AA      American Airlines Inc.
3 AS      Alaska Airlines Inc.
# ... with 13 more rows
```


Other related datasets

```
library(nycflights13)  
airlines  
flights  
planes  
airports  
weather
```

Digrammatic Representation of nycflights13 datasets




Figure 3: <https://r4ds.had.co.nz/relational-data.html>



Keys

Relational data

Keys

Joins

Visualising Geographic Data

Keys

Keys: variables used for connecting pairs of tables

- Primary key: variable/s which uniquely identify observations in their own table
 - weather: year, month, day, hour, origin
 - airports: faa
 - planes: tailnum
- Foreign key: variable/s which unique identify observations in another table
 - `flights$tailnum` is foreign key because it is a primary key in the table `planes`



Keys: Good practice

Good practice

- Identify primary key
- Check they uniquely identify observations
- If no primary key exists add a **surrogate key**

Example: Check key uniquely identifies observations 1

```
planes %>%  
  count(tailnum) %>%  
  filter(n > 1) # unique
```

```
# A tibble: 0 x 2
```

```
# ... with 2 variables: tailnum <chr>, n <int>
```

Example: Check key uniquely identifies observations 2

```
weather %>%  
  count(year, month, day, hour, origin) %>%  
  filter(n > 1) # not unique
```

A tibble: 3 x 6

	year	month	day	hour	origin	n
	<int>	<int>	<int>	<int>	<chr>	<int>
1	2013	11	3	1	EWR	2
2	2013	11	3	1	JFK	2
3	2013	11	3	1	LGA	2

Example: Adding a surrogate key 1

```
(tb <- tibble(x=c("A","B","B"), y=c(4,6,6)) )
```

```
# A tibble: 3 x 2
```

	x	y
	<chr>	<dbl>
1	A	4
2	B	6
3	B	6

```
tb %>% count(x, y)%>% filter(n > 1)
```

```
# A tibble: 1 x 3
```

	x	y	n
	<chr>	<dbl>	<int>
1	B	6	2

Example: Adding a surrogate key 2

```
tb %>% mutate(surrogate_key = row_number())
```

```
# A tibble: 3 x 3
```

	x	y	surrogate_key
	<chr>	<dbl>	<int>
1	A	4	1
2	B	6	2
3	B	6	3



Joins

Relational data

Keys

Joins

Visualising Geographic Data

Joins

- **Mutating joins:** add new variables from a data frame to another
- **Filtering joins:** filters a data frame based on whether they match another data frame

Relational data is usually stored in a **relational database management system (RDBMS)**.

Mutating Joins: Example

```
flights2 <- flights %>%  
  select(year:day, hour, origin, dest, tailnum, carrier)  
flights2
```

```
# A tibble: 336,776 x 8  
   year month   day hour origin dest  tailnum carrier  
   <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>  
1  2013     1     1     5 EWR    IAH   N14228  UA  
2  2013     1     1     5 LGA    IAH   N24211  UA  
3  2013     1     1     5 JFK    MIA   N619AA  AA  
# ... with 336,773 more rows
```

Suppose we want to add the airline name to the dataset.

Example

```
flights2 %>%  
  # remove columns for easier printing  
  select(-origin, -dest) %>%  
  # left join by carrier code  
  left_join(airlines, by = "carrier")
```

Types of mutating joins

- Inner joins – keeps only observations in x **and** y
- Outer joins
 - Left join – keeps all observations in x
 - Right join – keeps all observations in y
 - Full join – keeps all observations in x and y

Types of mutating joins




Figure 4: <https://r4ds.had.co.nz/relational-data.html>

Example

```
(x <- tribble( ~key, ~val_x, 1, "x1", 2, "x2", 3, "x3" ))  
(y <- tribble( ~key, ~val_y, 1, "y1", 2, "y2", 4, "y3" ))
```

```
# A tibble: 3 x 2
```

```
  key val_x
```

```
  <dbl> <chr>
```

```
1     1 x1
```

```
2     2 x2
```

```
3     3 x3
```

```
# A tibble: 3 x 2
```

```
  key val_y
```

```
  <dbl> <chr>
```

```
1     1 y1
```

```
2     2 y2
```

```
3     4 y3
```


Inner Joins

- New data has observations in **both** data sets.
- Unmatched observations are not included

```
x %>%  
inner_join(y, by="key")
```

```
# A tibble: 2 x 3  
  key val_x val_y  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2
```

Outer Joins: Left Join

```
x %>%  
  left_join(y, by="key")
```

```
# A tibble: 3 x 3  
  key val_x val_y  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2  
3     3 x3    <NA>
```

Outer Joins: Right join

```
x %>%  
  right_join(y, by="key")
```

```
# A tibble: 3 x 3  
  key val_x val_y  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2  
3     4 <NA>    y3
```

Outer Joins: Full join

```
x %>% full_join(y, by="key")
```

```
# A tibble: 4 x 3  
  key val_x val_y  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2  
3     3 x3    <NA>  
4     4 <NA>   y3
```

Ways of defining the key column

- Natural join: joins by all columns that appear in both tables x %>%
left_join(y, by = NULL)

```
flights2 %>% left_join(weather)
```

```
# A tibble: 336,776 x 18
```

	year	month	day	hour	origin	dest	tailnum	carrier	temp
	<int>	<int>	<int>	<dbl>	<chr>	<chr>	<chr>	<chr>	<dbl>
1	2013	1	1	5	EWB	IAH	N14228	UA	39.0
2	2013	1	1	5	LGA	IAH	N24211	UA	39.9
3	2013	1	1	5	JFK	MIA	N619AA	AA	39.0

```
# ... with 336,773 more rows, and 9 more variables:  
#   dewp <dbl>, humid <dbl>, wind_dir <dbl>,  
#   wind_speed <dbl>, wind_gust <dbl>, precip <dbl>,  
#   pressure <dbl>, visib <dbl>, time_hour <dtm>
```

Ways of defining the key column

- Specify key column using: by

```
flights2 %>% left_join(planes, by = "tailnum")
```

```
# A tibble: 336,776 x 16
```

```
  year.x month   day hour origin dest  tailnum carrier  
    <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>  
1  2013     1     1     5 EWR    IAH   N14228   UA  
2  2013     1     1     5 LGA    IAH   N24211   UA  
3  2013     1     1     5 JFK    MIA   N619AA   AA  
# ... with 336,773 more rows, and 8 more variables:  
#   year.y <int>, type <chr>, manufacturer <chr>,  
#   model <chr>, engines <int>, seats <int>, speed <int>,  
#   engine <chr>
```

Ways of defining the key column

- Same variable with different names in each table
 - `x %>% left_join(y, by=c("a" = "b"))`
 - Matches `x$a` with `y$b`

```
flights2 %>% left_join(airports, c("dest" = "faa"))
```

```
# A tibble: 336,776 x 15
```

```
  year month   day hour origin dest  tailnum carrier name
  <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>  <chr>
1  2013     1     1     5 EWR    IAH  N14228  UA      Geor~
2  2013     1     1     5 LGA    IAH  N24211  UA      Geor~
3  2013     1     1     5 JFK    MIA  N619AA  AA      Miam~
# ... with 336,773 more rows, and 6 more variables:
#   lat <dbl>, lon <dbl>, alt <dbl>, tz <dbl>, dst <chr>,
#   tzone <chr>
```

Filtering Joins

- `semi_join(x, y)` keep all in x that have match in y
- `anti_join(x, y)` drop all in x that have match in y

Similar to `filter`, but `anti-join` and `semi-join` scale better to use with more variables.

Filtering Joins: Example

```
(exams <- tribble(~studID, ~grade,  
                  1, "A",  
                  2, "B",  
                  3, "C+",  
                  5, "D" ))
```

A tibble: 4 x 2

studID grade

<dbl> <chr>

1	1	A
2	2	B
3	3	C+
4	5	D

Filtering Joins: Example

```
(classlist <- tribble( ~studentID, ~name,  
  1, "Charlotte",  
  2, "Zoe",  
  3, "Caitlin",  
  4, "Abel" ))
```

```
# A tibble: 4 x 2  
  studentID name  
    <dbl> <chr>  
1         1 Charlotte  
2         2 Zoe  
3         3 Caitlin  
4         4 Abel
```

Filtering Joins: Example

Questions of interest

1. Which students in the class sat the exam?
2. Which students didn't sit the exam?
3. Did any students not in the class sit the exam, what was their grade?

Filtering Joins: Semi-joins and Anti-joins

1. Which students in the class sat the exam?

```
classlist %>%  
  semi_join(exams, by=c("studentID"="studID"))
```

```
# A tibble: 3 x 2  
  studentID name  
    <dbl> <chr>  
1         1 Charlotte  
2         2 Zoe  
3         3 Caitlin
```

Filtering Joins: Semi-joins and Anti-joins

2. Which students didn't sit the exam?

```
classlist %>%  
  anti_join(exams, by=c("studentID"="studID"))
```

```
# A tibble: 1 x 2  
  studentID name  
    <dbl> <chr>  
1         4 Abel
```

Filtering Joins: Semi-joins and Anti-joins

3. Did any students not in the class sit the exam, what was their grade?

```
exams %>%  
  anti_join(classlist, by=c("studID"="studentID"))
```

```
# A tibble: 1 x 2  
  studID grade  
  <dbl> <chr>  
1      5 D
```

Example: Flights Data – Top 10 destinations

```
top_dest <- flights %>% count(dest, sort = TRUE) %>%  
  slice_head(n = 10) %>% print(n = 10)
```

```
# A tibble: 10 x 2
```

	dest	n
	<chr>	<int>
1	ORD	17283
2	ATL	17215
3	LAX	16174
4	BOS	15508
5	MCO	14082
6	CLT	14064
7	SFO	13331
8	FLL	12055
9	MIA	11728
10	DCA	9705

Filtering Joins: Semi-joins

1. Find all flights to the top destinations.

```
flights %>% semi_join(top_dest)
```

```
# A tibble: 141,145 x 19
```

	year	month	day	dep_time	sched~1	dep_d~2	arr_t~3	sched~4
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	542	540	2	923	850
2	2013	1	1	554	600	-6	812	837
3	2013	1	1	554	558	-4	740	728

```
# ... with 141,142 more rows, 11 more variables:
```

```
#   arr_delay <dbl>, carrier <chr>, flight <int>,
```

```
#   tailnum <chr>, origin <chr>, dest <chr>,
```

```
#   air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
#   minute <dbl>, time_hour <dtm>, and abbreviated
```

```
#   variable names 1: sched_dep_time, 2: dep_delay,
```

```
#   3: arr_time, 4: sched_arr_time
```


Filtering Joins: Semi-joins

1. Find all flights to the top destinations.

```
flights %>% semi_join(top_dest)
```

Equivalent to:

```
flights %>% filter(dest %in% top_dest$dest)
```

Filtering Joins: Anti-joins

2. Find flights whose plane isn't in planes

```
flights %>%  
  anti_join(planes, by = "tailnum") %>%  
  count(tailnum, sort = TRUE)
```

```
# A tibble: 722 x 2  
  tailnum      n  
  <chr>    <int>  
1 <NA>    2512  
2 N725MQ    575  
3 N722MQ    513  
# ... with 719 more rows
```

Filtering Joins: Anti-joins

2. Find flights whose plane isn't in planes

```
flights %>%  
  anti_join(planes, by = "tailnum") %>%  
  count(tailnum, sort = TRUE)
```

Equivalent to:

```
flights %>%  
  filter(! tailnum %in% planes$tailnum) %>%  
  count(tailnum, sort = TRUE)
```

Further topics in relational data

- Duplicate keys
- Join problems
- Set operations

If you are going to be working with relational data, you should read up about these topics. See Chapter 13 Wickham and Grolemund (2020).



Visualising Geographic Data

Relational data

Keys

Joins

Visualising Geographic Data

Join flights and airports

- Add longitude and latitude of airports to flights data

```
flights_loc <- flights %>%  
  select(carrier, flight, tailnum, origin, dest) %>%  
  inner_join(select(airports, faa, name, lat, lon),  
            by = c("origin"="faa")) %>%  
  inner_join(select(airports, faa, name, lat, lon),  
            by = c("dest"="faa"),  
            suffix = c("_origin", "_dest"))
```

Application: Join flights and airports

```
flights_loc
```

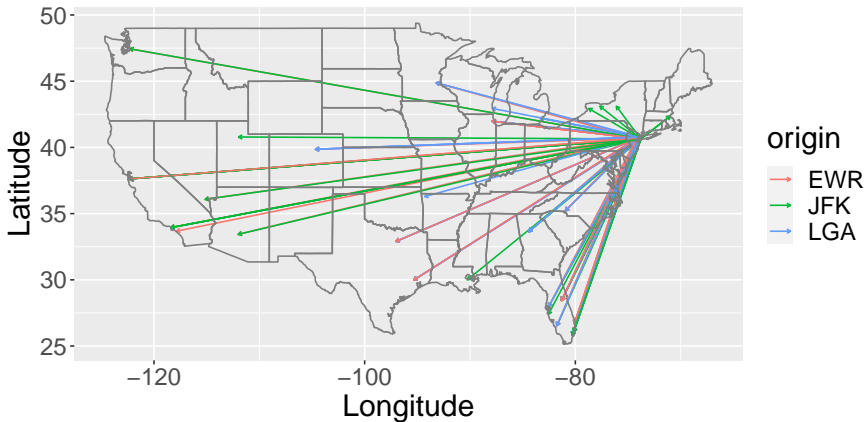
```
# A tibble: 329,174 x 11
  carrier flight tailnum origin dest  name_origin  lat_o~1
  <chr>      <int> <chr>    <chr> <chr> <chr>          <dbl>
1 UA          1545 N14228  EWR   IAH   Newark Libert~  40.7
2 UA          1714 N24211  LGA   IAH   La Guardia      40.8
3 AA          1141 N619AA  JFK   MIA   John F Kenned~  40.6
# ... with 329,171 more rows, 4 more variables:
#   lon_origin <dbl>, name_dest <chr>, lat_dest <dbl>,
#   lon_dest <dbl>, and abbreviated variable name
#   1: lat_origin
```

Visualise Flight Paths

```
flight_paths_plot <- flights_loc %>%  
  slice_head(n= 100) %>%  
  ggplot() +  
  geom_segment(mapping = aes(  
    x = lon_origin, xend = lon_dest,  
    y = lat_origin, yend = lat_dest,  
    col = origin),  
    arrow = arrow(length = unit(0.1, "cm")))) +  
  borders(database = "state") +  
  #borders(database = "world") +  
  coord_quickmap() +  
  labs(y = "Latitude", x = "Longitude")
```


Visualise Flight Paths (n = 100 flights)

```
flight_paths_plot
```



Most common destinations

```
(dest_freq <- flights %>%  
  count(dest) %>%  
  inner_join(airports, by=c("dest"="faa")) %>%  
  arrange(-n) )
```

A tibble: 101 x 9

	dest	n	name	lat	lon	alt	tz	dst	tzone
	<chr>	<int>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
1	ORD	17283	Chicago ~	42.0	-87.9	668	-6	A	Amer~
2	ATL	17215	Hartsfie~	33.6	-84.4	1026	-5	A	Amer~
3	LAX	16174	Los Ange~	33.9	-118.	126	-8	A	Amer~


... with 98 more rows

Plot of most common destinations from NYC airports

```
common_dest_plot <- ggplot() +  
  geom_point(data = dest_freq,  
            aes(x = lon , y = lat, size = n),  
            alpha = 0.5, col = "blue") +  
  borders(database = "state") +  
  #borders(database = "world") +  
  coord_fixed(1.3) +  
  guides(fill=FALSE, scale = "none")
```

Plot of most common destinations from NYC airports

`common_dest_plot`

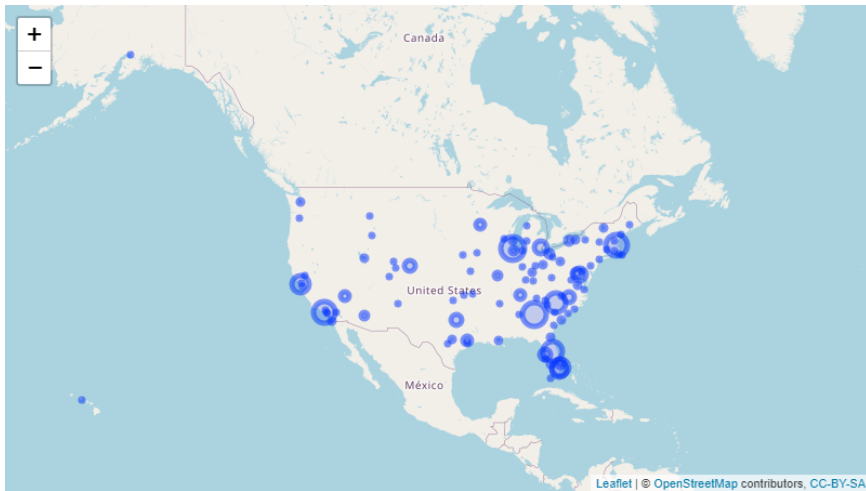


Leaflet plot

```
library(leaflet)
m <- leaflet(dest_freq) %>%
  addTiles() %>%
  #addProviderTiles(providers$Esri.WorldImagery) %>%
  addCircleMarkers(lng = ~lon, lat = ~lat,
                   popup = ~as.character(name),
                   label = ~as.character(name),
                   radius = ~n/1600) %>%
  setView(lng = -100, lat = 42, zoom = 3)

# save html widget as image
htmlwidgets::saveWidget(m, "temp.html", selfcontained=TRUE)
webshot2::webshot("temp.html",
                  file="Rfigs/leaflet_map.png",
                  cliprect="viewport",
                  vwidth = 800, vheight = 450)
```

Leaflet plot





Summary

Relational data

Keys

Joins

Visualising Geographic Data



Learning objectives

- Recognise relational data
- Understand the main types of mutating and filtering joins
- Join datasets using appropriate `tidyverse` join functions



References

Wickham, Hadley, and Garrett Grolemund. 2020. *R for Data Science: Import, Tidy, Transform, Visualize, and Model Data*.