

Lecture 7:

RNNs and Time Series Analysis

Yanbin Liu

Auckland University of Technology

April 18, 2024

- 1 Time Series Analysis
- 2 RNNs (LSTM, GRU)
- 3 Transformer Networks

Time Series Regression

- Time series regression is a statistical method for predicting a **future response** based on the **response history** (known as autoregressive dynamics) and the transfer of dynamics from relevant predictors.
- Time series regression can help us understand and predict the behavior of dynamic systems from experimental or observational data.
- Time series regression is employed for modeling and forecasting *economic, traffic, and weather* systems.

<https://au.mathworks.com/discovery/time-series-regression.html>

Linear Regression Model

$$\mathbf{y}_t = \mathbf{x}_t\beta + e_t$$

\mathbf{x}_t : includes *current* and *past* observations by time t

\mathbf{y}_t : an estimate of a *linear* relationship of the response

β : linear parameter estimates

e_t : innovation terms, difference between *observed* and *predicted*

Non-Linear Regression Model




$$\mathbf{y}_t = f(\mathbf{x}_t, e_t)$$

f can be:

- CNN
- RNN such as LSTM and GRU
- Transformer

Time Series Analysis

Deep Learning for Time Series Analysis




(a) The convolutional layer of a CNN with three groups (filters). (b) An artificial neural network (c) An LSTM block.

J. Gamboa. (2017) Deep Learning for Time-Series Analysis.


Time Series Analysis

MATLAB LSTM for Time Series Forecast




<https://au.mathworks.com/help/nnet/examples/time-series-forecasting-using-deep-learning.html>

MATLAB LSTM for Time Series Forecast



<https://au.mathworks.com/help/nnet/examples/time-series-forecasting-using-deep-learning.html>

MATLAB LSTM for Time Series Forecast with Updates



<https://au.mathworks.com/help/nnet/examples/time-series-forecasting-using-deep-learning.html>

Time Series Analysis


Questions?



Recurrent Neural Networks

Recurrent Neural Networks

RNNs are a family of neural networks for processing *sequential data*, which is a dynamical system. It uses the **same** *transition function* with the **same** *parameters* at every time step.




I. Goodfellow (2016) Deep Learning. MIT Press.

Recurrent Neural Networks

Unfolded RNNs

- RNNs produce an output at each time step and have **recurrent connections** between hidden units.
- RNNs read an entire sequence for processing and finally produce a single output.

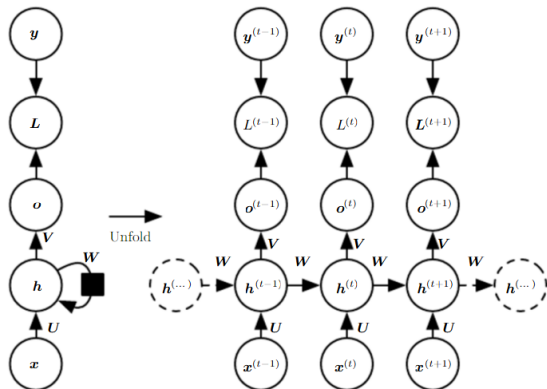


I. Goodfellow (2016) Deep Learning. MIT Press.


Recurrent Neural Networks

Unfolded RNNs

The notations are: Input x , state h , output o , loss function L , training target y , weights U, V , and W .




LSTM: Long Short-Term Memory




S. Hochreiter, et al. (1997) Long short-term memory, Neural computation, 9(8):1735-1780.


LSTM: Long Short-Term Memory



LSTM: Long Short-Term Memory




LSTM: Long Short-Term Memory



https://en.wikipedia.org/wiki/Long_short-term_memory.

LSTM: Long Short-Term Memory



$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$

Forget gate

$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$

Input gate

$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$

Output gate

$$\tilde{c}_t = \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$

Cell state

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

Cell Update

$$h_t = o_t \odot \sigma_h(c_t)$$

Output

https://en.wikipedia.org/wiki/Long_short-term_memory.


LSTM: Long Short-Term Memory

- Cell state c_t represents **long-term** memory
Hidden state h_t represents **short-term** memory
- Three **gates**: input gate, forget gate, and output gate.
- LSTM gates compute an activation, often using the logistic function.
- LSTM is well-suited to *classify*, *process* and *predict* time series given time lags of unknown size and duration between important events.
- LSTM was developed to deal with the *exploding* and *vanishing* gradient problems.

Web: https://en.wikipedia.org/wiki/Long_short-term_memory.

GRU: Gated Recurrent Unit

Only two gates: reset and update
 z_t and $1 - z_t$ updates.



r_t : reset gate


z_t : update gate

Web: https://en.wikipedia.org/wiki/Gated_recurrent_unit


GRU: Gated Recurrent Unit

Compare LSTM and GRU

- 3 gates vs. 2 gates
- More parameters vs. Fewer parameters
- Higher computation vs. lower computation
- Complex problems (intricate) vs. Smaller datasets (efficient)



LSTM




GRU

Questions?

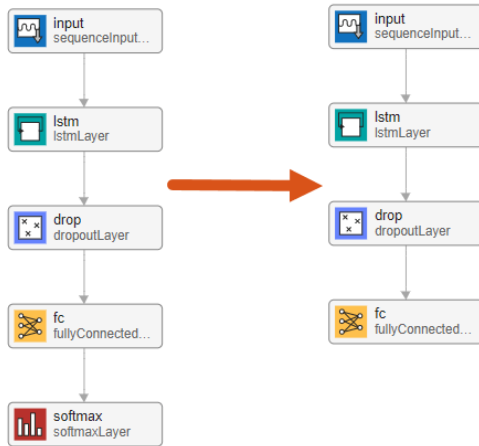


MATLAB LSTM Architecture

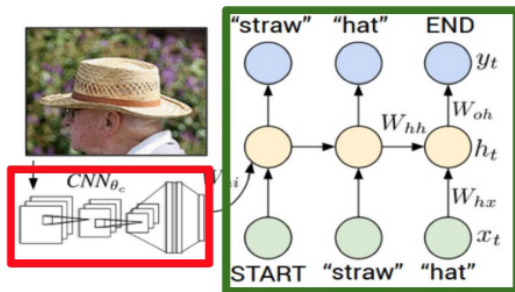


<https://au.mathworks.com/help/deeplearning/ug/long-short-term-memory-networks.html>

MATLAB LSTM in Time Series Analysis



ConvLSTM: Convolutional LSTM



Recurrent Neural Network

https://people.cs.pitt.edu/~kovashka/cs2770_sp17/vision_14_rnn.pdf

Rise of Transformer Model

- First designed for Natural Language Processing (NLP).
- Then extended to diverse fields: *computer vision, audio, video, etc.*
- Now the most powerful network architecture.
- Suitable for efficient and effective **time-series** prediction.

Microsoft PowerPoint

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

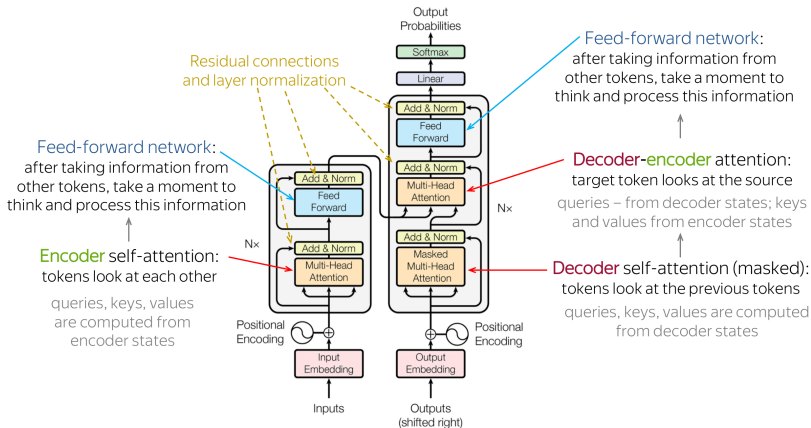
Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com


Transformer Networks

Transformer Architecture



Multi-head Attention

Scaled Dot-Product Attention




$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$


Transformer Networks

Multi-head Attention

Scaled Dot-Product Attention

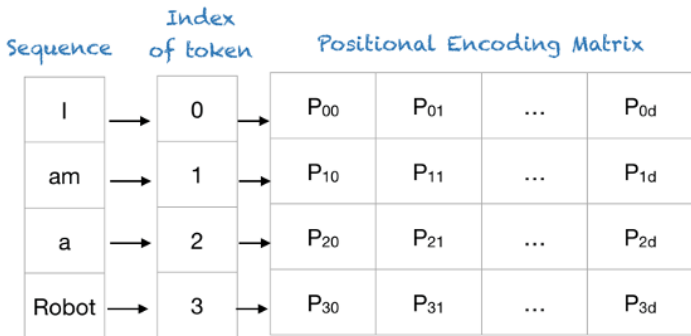


Multi-Head Attention



Transformer Networks

Position Encoding



Positional Encoding Matrix for the sequence 'I am a robot'

Transformer Networks

Position Encoding


Sequence	Index of token, k	Positional Encoding Matrix with $d=4$, $n=100$			
		$i=0$	$i=0$	$i=1$	$i=1$
I	0	$P_{00}=\sin(0)$ = 0	$P_{01}=\cos(0)$ = 1	$P_{02}=\sin(0)$ = 0	$P_{03}=\cos(0)$ = 1
am	1	$P_{10}=\sin(1/1)$ = 0.84	$P_{11}=\cos(1/1)$ = 0.54	$P_{12}=\sin(1/10)$ = 0.10	$P_{13}=\cos(1/10)$ = 1.0
a	2	$P_{20}=\sin(2/1)$ = 0.91	$P_{21}=\cos(2/1)$ = -0.42	$P_{22}=\sin(2/10)$ = 0.20	$P_{23}=\cos(2/10)$ = 0.98
Robot	3	$P_{30}=\sin(3/1)$ = 0.14	$P_{31}=\cos(3/1)$ = -0.99	$P_{32}=\sin(3/10)$ = 0.30	$P_{33}=\cos(3/10)$ = 0.96

Positional Encoding Matrix for the sequence 'I am a robot'

Transformer Networks

Mask Attention Mechanism


- Predicting y_t can only use x_1, \dots, x_{t-1} .
- The model can only *use past rather than future information*.
- For example, you cannot use tomorrow's weather/stock to predict today or yesterday.



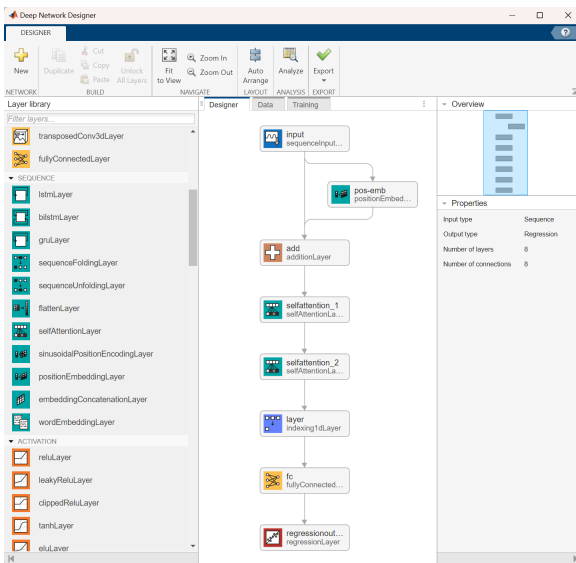
raw attention weights



mask



Transformer Networks in MATLAB



<https://au.mathworks.com/fileexchange/161016-transformer-networks-for-time-series-prediction>

Transformer Networks

Questions?

